

```
1 import torch
```

3.1 Basics of Autograd

```
1 # Creating a tensor with requires_grad=True
2 x = torch.tensor([2.0, 3.0], requires_grad=True)
3 print(f"Tensor x with requires_grad=True: {x}")
4
5 # Performing operations
6 y = x + 2
7 print(f"y = x + 2: {y}")
8
9 # More operations
10 z = y * y * 3
11 out = z.mean()
12 print(f"z = y * y * 3: {z}")
13 print(f"out = z.mean(): {out}")
14
```

→ Tensor x with requires_grad=True: tensor([2., 3.], requires_grad=True)
y = x + 2: tensor([4., 5.], grad_fn=<AddBackward0>)
z = y * y * 3: tensor([48., 75.], grad_fn=<MulBackward0>)
out = z.mean(): 61.5

3.2 Computing Gradients

```
1 # Backpropagation
2 out.backward()
3 print(f"Gradients of x: {x.grad}")
4
5 # Note: The gradient will be calculated for the scalar output 'out' w.r.t. the input 'x'
6
```

→ Gradients of x: tensor([12., 15.])

3.3 Stopping Gradient Tracking

```
1 # Using torch.no_grad()
2 with torch.no_grad():
3     y = x + 2
4     print(f"y = x + 2 with no_grad: {y}")
5
6 # Using detach() method
7 y = x.detach()
8 print(f"y = x.detach(): {y}")
9
```


→ y = x + 2 with no_grad: tensor([4., 5.])
y = x.detach(): tensor([2., 3.])

3.4 Gradient Accumulation

```

1 # Create a tensor with requires_grad=True
2 x = torch.tensor([1.0, 2.0, 3.0], requires_grad=True)
3
4 # Perform first set of operations
5 y1 = x + 2
6 z1 = y1 * y1
7 z1 = z1.mean()
8
9 # Compute first set of gradients
10 z1.backward()
11 print(f"Gradients of x after first backward pass: {x.grad}")
12
13 # Perform second set of operations without resetting gradients
14 y2 = x * 3
15 z2 = y2.mean()
16
17 # Compute second set of gradients
18 z2.backward()
19 print(f"Gradients of x after second backward pass (accumulated): {x.grad}")
20
21 # Note: Gradients are accumulated in x.grad
22 # Reset gradients for a clean state
23 x.grad.zero_()
24
25 # Perform third set of operations
26 y3 = x * 4
27 z3 = y3.mean()
28
29 # Compute third set of gradients
30 z3.backward()
31 print(f"Gradients of x after third backward pass (after zeroing): {x.grad}")

```


 Gradients of x after first backward pass: tensor([2.0000, 2.6667, 3.3333])
 Gradients of x after second backward pass (accumulated): tensor([3.0000, 3.6667, 4.3333])
 Gradients of x after third backward pass (after zeroing): tensor([1.3333, 1.3333, 1.3333])

✓ 3.5 Custom Gradients with Function

```

1 class MyReLU(torch.autograd.Function):
2     @staticmethod
3     def forward(ctx, input):
4         """
5         The forward method computes the output of the function given the input.
6         ctx is a context object that can be used to stash information for backward computation.
7         input is the input tensor.
8         """
9         ctx.save_for_backward(input) # Save input tensor for backward pass
10        return input.clamp(min=0) # Apply ReLU (Rectified Linear Unit) operation
11
12    @staticmethod
13    def backward(ctx, grad_output):
14        """
15        The backward method computes the gradient of the function w.r.t. its input.
16        ctx is a context object that contains information saved during the forward pass.
17        grad_output is the gradient of the loss w.r.t. the output of this function.
18        """
19        input, = ctx.saved_tensors # Retrieve saved input tensor
20        grad_input = grad_output.clone() # Clone the gradient output tensor
21        grad_input[input < 0] = 0 # Apply ReLU derivative: gradient is zero where input was negative
22        return grad_input # Return the gradient w.r.t. the input
23
24
25 # Create a tensor with requires_grad=True
26 x = torch.tensor([-2.0, 0.0, 3.0], requires_grad=True)
27
28 # Apply the custom ReLU function
29 relu = MyReLU.apply
30 y = relu(x)
31
32 # Perform backpropagation
33 y.backward(torch.tensor([1.0, 1.0, 1.0]))
34 print(f"Custom ReLU function gradients: {x.grad}")
35

```

 Custom ReLU function gradients: tensor([0., 1., 1.])

✓ Exercises

1. Implement a Custom Sigmoid Function: Create a custom autograd function for the sigmoid activation function and verify its gradients.
2. Custom Function for Squaring: Implement a custom autograd function that squares its input and compute its gradients.
3. Gradient Checking: Implement numerical gradient checking to verify the correctness of custom gradients.

```
1 import torch
2 from torch.autograd import Function
3
4 # Custom sigmoid function
5 class CustomSigmoid(Function):
6     @staticmethod
7     def forward(ctx, input):
8         sigmoid_output = 1 / (1 + torch.exp(-input))
9         ctx.save_for_backward(sigmoid_output)
10        return sigmoid_output
11
12    @staticmethod
13    def backward(ctx, grad_output):
14        sigmoid_output, = ctx.saved_tensors
15        grad_input = grad_output * sigmoid_output * (1 - sigmoid_output)
16        return grad_input
17
18 # Usage example
19 if __name__ == "__main__":
20     # Input tensor
21     x = torch.tensor([0.5, -1.0, 2.0], requires_grad=True)
22
23     # Apply custom sigmoid
24     custom_sigmoid = CustomSigmoid.apply
25     y = custom_sigmoid(x)
26
27     # Print the forward result
28     print("Forward result:", y)
29
30     # Compute gradients
31     y.sum().backward()
32
33     # Print the gradients
34     print("Gradients:", x.grad)
35
```