

4.1 Defining a Neural Network

```

1 import torch
2 import torch.nn as nn
3
4 # Define a simple neural network
5 class SimpleNN(nn.Module):
6     def __init__(self):
7         super(SimpleNN, self).__init__()
8         # Define layers
9         self.fc1 = nn.Linear(3, 5) # Input layer to hidden layer
10        self.fc2 = nn.Linear(5, 1) # Hidden layer to output layer
11
12    def forward(self, x):
13        # Define forward pass
14        x = torch.relu(self.fc1(x)) # Apply ReLU activation
15        x = self.fc2(x)
16        return x
17
18 # Instantiate the network
19 net = SimpleNN()
20 print(net)
21

```

```

SimpleNN(
  (fc1): Linear(in_features=3, out_features=5, bias=True)
  (fc2): Linear(in_features=5, out_features=1, bias=True)
)

```

4.2 Loss Function

```

1 # Define Mean Squared Error loss function
2 criterion = nn.MSELoss()

```

4.3 Optimizer

```

1 import torch.optim as optim
2
3 # Define an optimizer
4 optimizer = optim.SGD(net.parameters(), lr=0.01)
5

```

4.4 Preparing Data

```

1 # Example training data
2 inputs = torch.tensor([[1.0, 2.0, 3.0], [4.0, 5.0, 6.0], [7.0, 8.0, 9.0]], dtype=torch.float32)
3 targets = torch.tensor([[1.0], [2.0], [3.0]], dtype=torch.float32)
4
5 # Create a DataLoader
6 from torch.utils.data import DataLoader, TensorDataset
7
8 dataset = TensorDataset(inputs, targets)
9 data_loader = DataLoader(dataset, batch_size=2, shuffle=True)
10

```

4.5 Training Loop

```

1 # Training loop
2 num_epochs = 100
3 for epoch in range(num_epochs):
4     for batch_inputs, batch_targets in data_loader:
5         # Zero the gradient buffers
6         optimizer.zero_grad()
7
8         # Forward pass
9         outputs = net(batch_inputs)
10
11        # Compute loss
12        loss = criterion(outputs, batch_targets)
13
14        # Backward pass
15        loss.backward()
16
17        # Update weights
18        optimizer.step()
19
20    if epoch % 10 == 0:

```

4.6 Evaluating the Model

```

21 print("\nTesting results:")
22
23 # Example test data
24 test_inputs = torch.tensor([[10.0, 11.0, 12.0]], dtype=torch.float32)
25 test_targets = torch.tensor([[4.0]], dtype=torch.float32)
26
27 # Forward pass on test data
28 with torch.no_grad():
29     test_outputs = net(test_inputs)
30     test_loss = criterion(test_outputs, test_targets)
31
32 print(f"Test Loss: {test_loss.item()}")
33

```

Test Loss: 2.29196302825585e-05

Exercises

1. Define a More Complex Network: Create a deeper network with more layers and different activation functions. Experiment with different layer sizes and configurations.
2. Experiment with Different Loss Functions: Try using a different loss function, such as Cross Entropy Loss, and modify the network for a classification task.
3. Optimize with Different Optimizers: Compare the performance of different optimizers (e.g., SGD vs. Adam) on the same task. Adjust learning rates and observe the effects.
4. Add Regularization: Implement L2 regularization by adding a penalty to the loss function. Experiment with dropout layers to prevent overfitting.
5. Data Augmentation: Use data augmentation techniques to artificially increase the size of the training set. Implement these techniques and observe their effects on model performance.